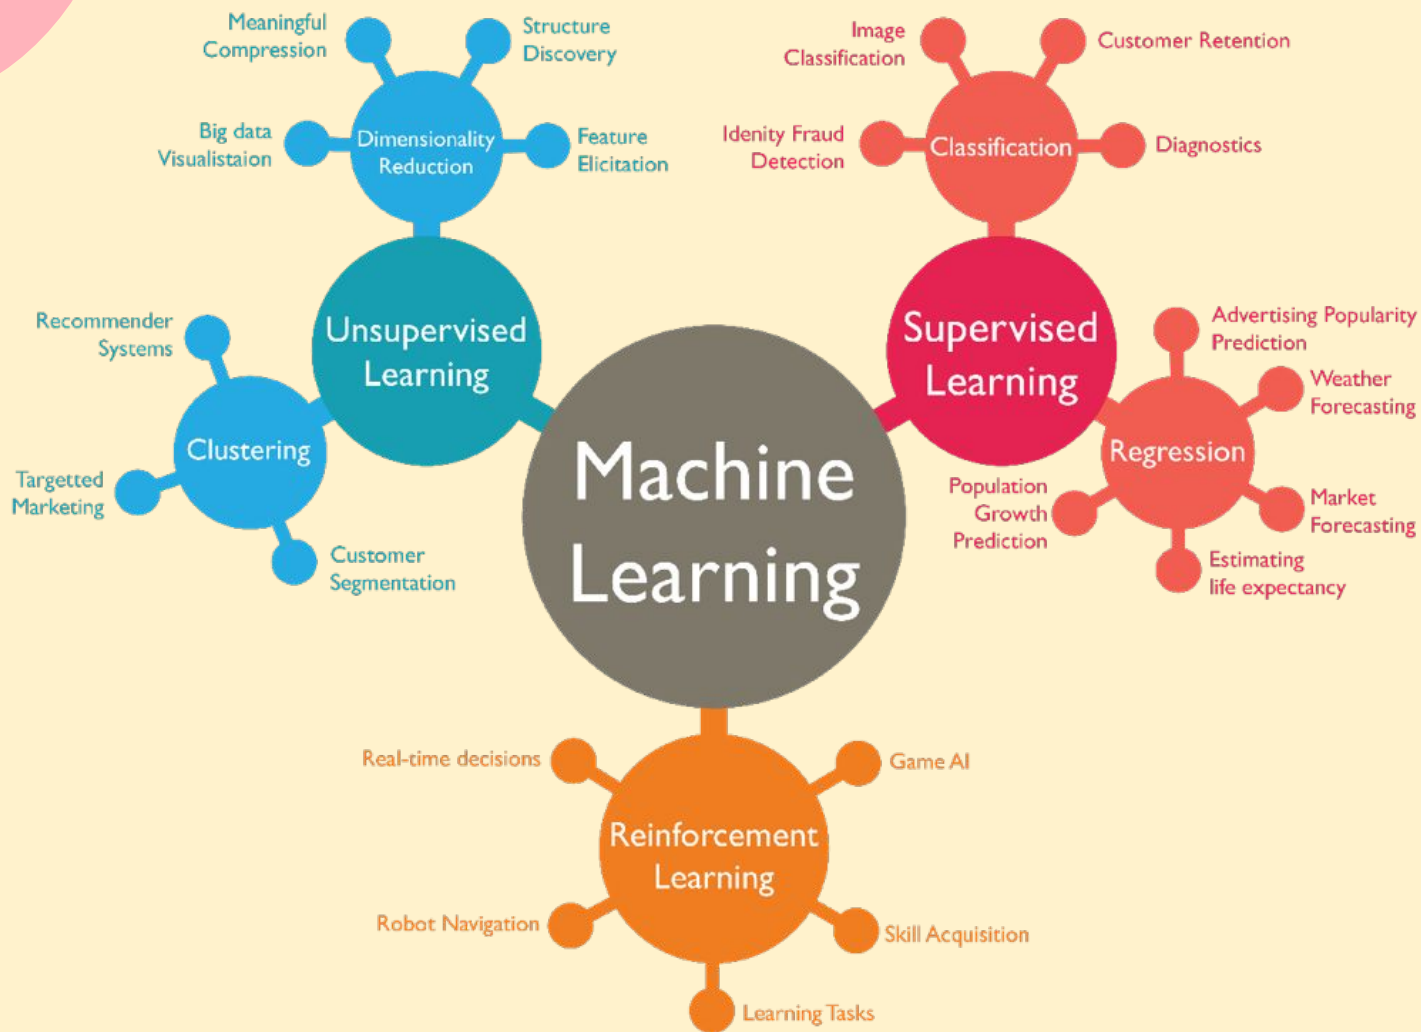


Capítol 1.2: Machine Learning


Aina Palacios



How Supervised Machine Learning Works


STEP 1

Provide the machine learning algorithm categorized or "labeled" input and output data from to learn




STEP 2

Feed the machine new, unlabeled information to see if it tags new data appropriately. If not, continue refining the algorithm




TYPES OF PROBLEMS TO WHICH IT'S SUITED



CLASSIFICATION

Sorting items into categories




REGRESSION

Identifying real values (dollars, weight, etc.)

How **Unsupervised** Machine Learning Works


STEP 1

Provide the machine learning algorithm uncategorized, unlabeled input data to see what patterns it finds




STEP 2

Observe and learn from the patterns the machine identifies




TYPES OF PROBLEMS TO WHICH IT'S SUITED



CLUSTERING

Identifying similarities in groups

For Example: Are there patterns in the data to indicate certain patients will respond better to this treatment than others?



ANOMALY DETECTION

Identifying abnormalities in data

For Example: Is a hacker intruding in our network?

Entorn de treball

Pandas

Per poder treballar la
base de dades



Llibreria Sklearn

Per poder aplicar
Machine Learning així
com preprocessats



Com aplicar Machine Learning

1. Determinar objectiu i estudi de les dades
2. Preprocessat
3. Procés enginyeria
4. Determinar train i test
5. Escollir i entrenar el model
6. Mètriques i validació creuada
7. Modificar paràmetres

Com aplicar Machine Learning

1. **Determinar objectiu i estudi de les dades**
2. Preprocessat
3. Procés enginyeria
4. Determinar train i test
5. Escollir i entrenar el model
6. Mètriques i validació creuada
7. Modificar paràmetres

1. Determinar Objectiu i estudi de les dades

Primerament, hem de determinar quin és l'objectiu que estem perseguint. A vegades hem de fixar-nos que les mostres recollides contenen informació a priori i posterior de l'esdeveniment:

- Per exemple, si busquem si un paquet arribarà tard al seu destí i tenim una base de dades amb tot el recorregut del paquet, la informació a posteriori de l'entrega, així com els motius de retard, són dades que no podem utilitzar, ja que no són indicadors prioris del retard de l'entrega.

El nostre objectiu pot ser classificar, predir una regressió o aplicar una tècnica de clústering.

Per poder entendre millor el que estem fent, és important fer un estudi de totes les característiques de la meva mostra. D'aquesta manera, podem aplicar **preprocessats** de les mostres per ajudar als algoritmes a treballar millor.

Anàlisi descriptiu i gràfic

Com aplicar Machine Learning

1. Determinar objectiu i estudi de les dades
2. **Preprocessat**
3. Procés enginyeria
4. Determinar train i test
5. Escollir i entrenar el model
6. Mètriques i validació creuada
7. Modificar paràmetres

2. Preprocessat

És important netejar les dades abans d'aplicar cap algoritme sobre elles

1. **Dades no existents:** Nans, Unknown, Null, ext.
 - a. Si són numèriques podem posar a 0, fer interpolació, mitjana,...
 - b. Si són categòriques podem posar un identificador
 - c. Si no es poden substituir, hauriem d'eliminar les dades
2. **Dades anòmales:** Hi ha dades que poden no tenir sentit, aplicar el coneixement previ. Un outlier no és una dada anòmala.

Com aplicar Machine Learning



1. Determinar objectiu i estudi de les dades
2. Preprocessat
- 3. Procés enginyeria**
4. Determinar train i test
5. Escollir i entrenar el model
6. Mètriques i validació creuada
7. Modificar paràmetres

3. Procés enginyeria

No aplicar procés sobre el target o invertir el procés

3. **Procés d'enginyeria:** Existeixen diferents transformacions que ajuden al model a predir millor.

- a. Si són gaussianes -> Estandardització
- b. Si la seva distribució no és normal -> Normalització
- c. Si conté outliers -> RobustScaler o altres
- d. Si són categòriques -> Dummys o enumeració
- e. Altres: Polimorfisme, transformacions, cícliques,...




Com aplicar Machine Learning

1. Determinar objectiu i estudi de les dades
2. Preprocessat
3. Procés enginyeria
4. **Determinar train i test**
5. Escollir i entrenar el model
6. Mètriques i validació creuada
7. Modificar paràmetres

4. Determinar train i test

Si treballem amb Supervised Machine Learning, necessitem dividir les dades entre train i test per tal de poder determinar el millor **model** per aconseguir el nostre objectiu.



El test l'utilitzarem per veure com de bé el model funciona. Normalment representa un 30% o 20% de les dades.

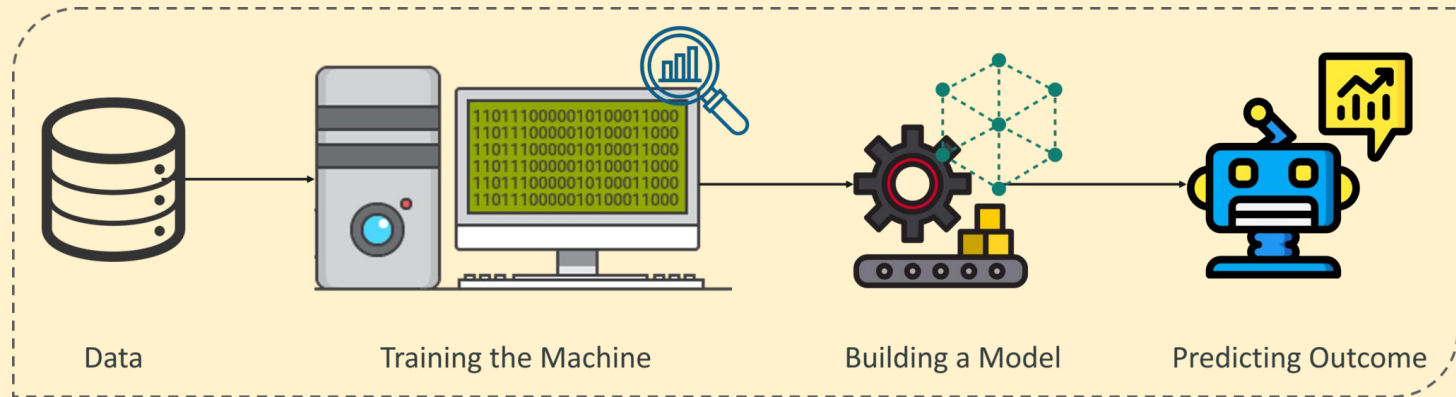
```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2)
```

Com aplicar Machine Learning

1. Determinar objectiu i estudi de les dades
2. Preprocessat
3. Procés enginyeria
4. Determinar train i test
- 5. Escollir i entrenar el model**
6. Mètriques i validació creuada
7. Modificar paràmetres

Machine Learning

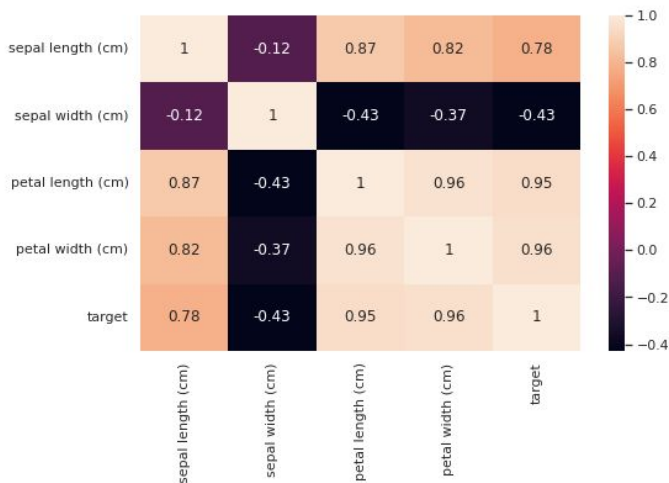
- Enllaç a curs de [Machine Learning](#)
- Enllaç Machine Learning amb [Sklearn](#)
- [Cheating sheet](#) Machine Learning amb Sklearn



Classification

```
corr = df.corr()  
sns.heatmap(corr, annot=True)
```

<AxesSubplot:>



```
from sklearn import neighbors, datasets, preprocessing  
from sklearn.model_selection import train_test_split  
from sklearn.metrics import accuracy_score
```

```
iris = datasets.load_iris()
```

```
import pandas as panda  
import numpy as np
```

```
df= pd.DataFrame(data= np.c_[iris['data'], iris['target']],  
                  columns= iris['feature_names'] + ['target'])  
df.head()
```

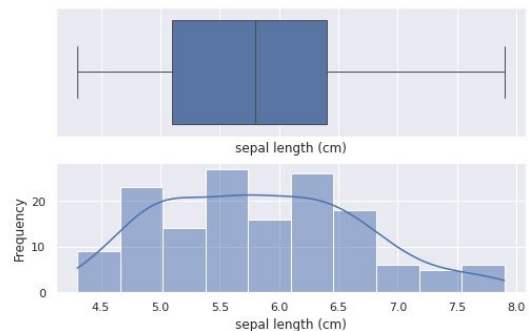
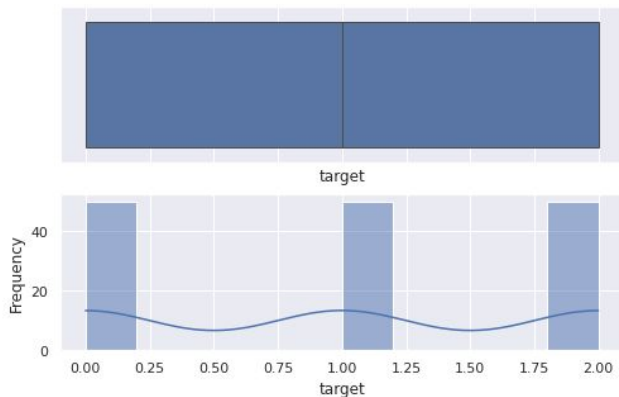
	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	target
0	5.1	3.5	1.4	0.2	0.0
1	4.9	3.0	1.4	0.2	0.0
2	4.7	3.2	1.3	0.2	0.0
3	4.6	3.1	1.5	0.2	0.0
4	5.0	3.6	1.4	0.2	0.0

Classification


```
import matplotlib.pyplot as plt
import seaborn as sns

import warnings
warnings.filterwarnings("ignore")

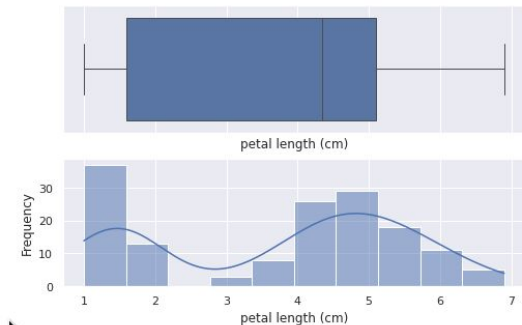
for i in df.columns:
    plt.figure()
    plt.tight_layout()
    sns.set(rc={"figure.figsize":(8, 5)})
    f, (ax_box, ax_hist) = plt.subplots(2, sharex=True)
    plt.gca().set(xlabel= i, ylabel='Frequency')
    sns.boxplot(df[i], ax=ax_box, linewidth= 1.0)
    sns.histplot(df[i], ax=ax_hist, bins = 10, kde=True)
```




<Figure size 576x360 with 0 Axes>



<Figure size 576x360 with 0 Axes>



<Figure size 576x360 with 0 Axes>





1.2 Preprocessat



```
df.isnull().values.any()
```

```
False
```

No tenim valors nuls i considerem que no existeixen dades anomeles



<Figure size 576x360 with 0 Axes>



- La distribució de sepal width és normal, per tant aplicarem estandarització
- Les altres variables no contenen outliers, per tant utilitzarem normalització

```
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import MinMaxScaler
```

```
standColumns = ['sepal width (cm)']
scalerStand = preprocessing.StandardScaler().fit(df[standColumns])
df[standColumns] = scalerStand.transform(df[standColumns])
```

```
normColumns = ['sepal length (cm)', 'petal length (cm)', 'petal width (cm)']
scalerNorm = preprocessing.MinMaxScaler().fit(df[normColumns])
df[normColumns] = scalerNorm.transform(df[normColumns])
```

```
import matplotlib.pyplot as plt
import seaborn as sns
```

```
import warnings
warnings.filterwarnings("ignore")
```


```
for i in df.columns:
    plt.figure()
    plt.tight_layout()
    sns.set(rc={"figure.figsize":(8, 5)})
    f, (ax_box, ax_hist) = plt.subplots(2, sharex=True)
    plt.gca().set(xlabel= i, ylabel='Frequency')
    sns.boxplot(df[i], ax=ax_box, linewidth= 1.0)
    sns.histplot(df[i], ax=ax_hist, bins= 10, kde=True)
```

1.3 Test/train

```
from sklearn.model_selection import train_test_split  
  
X = df.drop(['target'],axis=1)  
y = df['target']  
  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2)
```

Hem utilitzat un test del 20% perquè tenim poques mostres

Classification



1.4 Models

```
from sklearn.linear_model import LogisticRegression
lr = LogisticRegression(random_state=42)
```

```
from sklearn.svm import SVC
svc = SVC(kernel='linear')
```

```
from sklearn.ensemble import RandomForestClassifier
rf = RandomForestClassifier(max_depth=2, random_state=42)
```

Training the models

```
lr.fit(X_train, y_train)
svc.fit(X_train, y_train)
rf.fit(X_train, y_train)
```

...

```
y_pred_lr = lr.predict(X_test)
y_pred_svc = svc.predict(X_test)
y_pred_rf = rf.predict(X_test)
```

```
y_pred_rf
```

```
array([2, 1, 2, 0, 2, 1, 2, 0, 0, 1, 1, 1, 1, 0, 0, 2, 0, 0, 0, 2, 2, 1,
       1, 2, 2, 2, 1, 2, 1, 0])
```

Com aplicar Machine Learning


1. Determinar objectiu i estudi de les dades
2. Preprocessat
3. Procés enginyeria
4. Determinar train i test
5. Escollir i entrenar el model
6. **Mètriques i validació creuada**
7. Modificar paràmetres

Classification

Model Evaluation



Confusion Matrix

`confusion_matrix(y_test, y_pred)`



ROC & AUC

`metrics.auc(fpr, tpr)`



How many retrieved items are relevant?

$$\text{Precision} = \frac{\text{true positives}}{\text{true positives} + \text{false positives}}$$

How many relevant items are retrieved?

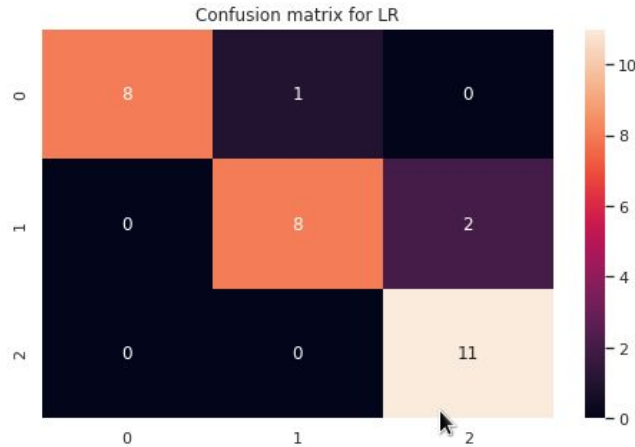
$$\text{Recall} = \frac{\text{true positives}}{\text{true positives} + \text{false negatives}}$$

5 Evaluació

```
from sklearn.metrics import confusion_matrix
cf_matrix_lr = confusion_matrix(y_test, y_pred_lr)
cf_matrix_svc = confusion_matrix(y_test, y_pred_svc)
cf_matrix_rf = confusion_matrix(y_test, y_pred_rf)

sns.heatmap(cf_matrix_lr, annot=True).set(title='Confusion matrix for LR')


[Text(0.5, 1.0, 'Confusion matrix for LR')]
```




```
from sklearn.metrics import f1_score
f1_lr = f1_score(y_test, y_pred_lr, average='macro')
f1_svc = f1_score(y_test, y_pred_svc, average='macro')
f1_rf = f1_score(y_test, y_pred_rf, average='macro')
print("F1 for LR: {:.2f}, SVC: {:.2f}, RF: {:.2f}".format(f1_lr, f1_svc, f1_rf))
```

F1 for LR: 0.90, SVC: 1.00, RF: 1.00

Confusion matrix for SVC



Confusion matrix for RF



Validació creuada

Quan creem el test/train pot donar la casualitat que el nostre set afavoreix el nostre model. Per tal de trobar el valor real dels nostres paràmetres, utilitzem CrossValidation.



Cross Validation

```
from sklearn.model_selection import cross_val_score
print(cross_val_score(lr, X, y, cv=5, scoring='f1_macro'))
```

Validació creuada

```
cv_lr = cross_val_score(lr, X, y, cv=5, scoring='f1_macro')
print("F1 for LR mean: {:.2f}, std: {:.2f}".format(cv_lr.mean(), cv_lr.std()) )

cv_svc = cross_val_score(svc, X, y, cv=5, scoring='f1_macro')
print("F1 for SVC mean: {:.2f}, std: {:.2f}".format(cv_svc.mean(), cv_svc.std()) )

cv_rf = cross_val_score(rf, X, y, cv=5, scoring='f1_macro')
print("F1 for RF mean: {:.2f}, std: {:.2f}".format(cv_rf.mean(), cv_rf.std()) )
```

```
F1 for LR mean: 0.92, std: 0.05
F1 for SVC mean: 0.96, std: 0.03
F1 for RF mean: 0.95, std: 0.03
```

Com aplicar Machine Learning



1. Determinar objectiu i estudi de les dades
2. Preprocessat
3. Procés enginyeria
4. Determinar train i test
5. Escollir i entrenar el model
6. Mètriques i validació creuada
7. **Modificar paràmetres**

7. Modificar paràmetres

El model conté diferents **paràmetres** que modifiquen el seu comportament.

És important intentar trobar aquells paràmetres que donen els millors resultats.

Podem intentar trobar-los directament o utilitzant el **GridSearch** o un **Randomized Parameter Optimization**, que ens provarà totes les combinacions possibles que li diguem i trobarà el millor resultat.



6. Millors paràmetres

```
from sklearn.model_selection import RandomizedSearchCV


params = {"solver": ['newton-cg', 'lbfgs', 'liblinear', 'sag', 'saga'], "penalty": ["l1", "l2", "elasticnet"]}
rsearch = RandomizedSearchCV(estimator=lr,
                             param_distributions=params, cv=5,
                             n_iter=14, random_state=42, scoring='f1_macro')

rsearch.fit(X, y)
print(rsearch.best_score_)
print(rsearch.best_params_)
```


```
0.9530138477506899
{'solver': 'saga', 'penalty': 'l1'}
```

```
y_pred_RS_ls = LogisticRegression(random_state=42, solver = 'saga', penalty = 'l1')\
    .fit(X_train, y_train).predict(X_test)
cf_matrix_RS_lr = confusion_matrix(y_test, y_pred_RS_ls)
sns.heatmap(cf_matrix_RS_lr, annot=True).set(title='Confusion matrix for LR')
```

```
[Text(0.5, 1.0, 'Confusion matrix for LR')]
```



Regression



Machine Learning

Regression

Utilitzarem la base de dades Iris, que ens proporciona un Data Set sobre tipus de plantes i els seus pètals

```
from sklearn import neighbors, datasets, preprocessing
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
```

```
diabetes = datasets.load_diabetes()
```


```
import pandas as panda
import numpy as np
```

```
df = panda.DataFrame(diabetes.data, columns=diabetes.feature_names)
df['Progression'] = diabetes.target
```


```
df.head()
```

	age	sex	bmi	bp	s1	s2	s3	s4	s5	s6	Progression
0	0.038076	0.050680	0.061696	0.021872	-0.044223	-0.034821	-0.043401	-0.002592	0.019907	-0.017646	151.0
1	-0.001882	-0.044642	-0.051474	-0.026328	-0.008449	-0.019163	0.074412	-0.039493	-0.068332	-0.092204	75.0
2	0.085299	0.050680	0.044451	-0.005670	-0.045599	-0.034194	-0.032356	-0.002592	0.002861	-0.025930	141.0
3	-0.089063	-0.044642	-0.011595	-0.036656	0.012191	0.024991	-0.036038	0.034309	0.022688	-0.009362	206.0
4	0.005383	-0.044642	-0.036385	0.021872	0.003935	0.015596	0.008142	-0.002592	-0.031988	-0.046641	135.0


Regression





age




<Figure size 576x360 with 0 Axes>




sex




bmi




<Figure size 576x360 with 0 Axes>





bp







<Figure size 576x360 with 0 Axes>




<Figure size 576x360 with 0 Axes>




Regression





s5




<Figure size 576x360 with 0 Axes>



s6



Progression




3 Test/train

```
from sklearn.model_selection import train_test_split

X = df.drop(['Progression'],axis=1)
y = df['Progression']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3)
```



4 Models

```
from sklearn.linear_model import LinearRegression
lr = LinearRegression()
```

```
from sklearn.svm import SVR
svr = SVR()
```

Training the models

```
lr.fit(X_train, y_train)
svr.fit(X_train, y_train)
```

...

```
y_pred_lr = lr.predict(X_test)
y_pred_svr = svr.predict(X_test)
```

```
y_pred_lr[1:5]
```

```
array([135.69448628, 196.53507727, 74.45375974, 87.28375355])
```

5 Evaluació

```
from sklearn.metrics import r2_score
```

```
r2_lr = r2_score(y_test, y_pred_lr)
r2_svr = r2_score(y_test, y_pred_svr)
```

```
from sklearn.metrics import mean_squared_error
```

```
mse_lr = mean_squared_error(y_test, y_pred_lr)
mse_svr = mean_squared_error(y_test, y_pred_svr)
```

```
metrics = panda.DataFrame( data = [['Linear Regression', r2_lr, mse_lr],
                                   ['Support Vector Machines', r2_svr, mse_svr]],
                           columns = ['Model', 'R2', 'MSE'])
metrics.head()
```

	Model	R2	MSE
0	Linear Regression	0.502632	3289.974895
1	Support Vector Machines	0.157471	5573.138008

Cross Validation

```
from sklearn.model_selection import cross_val_score
```

```
lr = LinearRegression()
cv_lr = cross_val_score(lr, X, y, cv=5, scoring='r2')
cv_svr = cross_val_score(svr, X, y, cv=5, scoring='r2')
```

```
cv_lr_mean = [cv_lr.mean(), cv_svr.mean()]
metrics['R2 with CV'] = cv_lr_mean
metrics.head()
```

	Model	R2	MSE	R2 with CV
0	Linear Regression	0.255382	3761.752362	0.482316
1	Support Vector Machines	0.126417	4413.269770	0.161335

6. Millors paràmetres

```
from sklearn.model_selection import RandomizedSearchCV
param = {'kernel' : ('linear', 'poly', 'rbf', 'sigmoid'), 'C' : [1,5,10], 'degree' : [3,8],
         'coef0' : [0.01,10,0.5], 'gamma' : ('auto', 'scale')}
```

```
lr = LinearRegression()
rsearch = RandomizedSearchCV(estimator=svr,
                             param_distributions=param, cv=5,
                             n_iter=10, random_state=42, scoring='r2')

rsearch.fit(X, y)
print(rsearch.best_score_)
print(rsearch.best_params_)
```

```
0.47480829897832244
{'kernel': 'linear', 'gamma': 'scale', 'degree': 3, 'coef0': 0.5, 'C': 10}
```

```
svr_bp = SVR(kernel='linear', gamma='scale', degree= 3, coef0= 0.5, C= 10).fit(X_train, y_train)
y_pred_svr_bp = svr_bp.predict(X_test)
```

```
metrics_bp = panda.DataFrame( data = [['SVR best params', r2_score(y_test, y_pred_svr_bp),
                                     mean_squared_error(y_test, y_pred_svr_bp),
                                     cross_val_score(svr_bp, X, y, cv=5, scoring='r2').mean()]],
                              columns = ['Model', 'R2', 'MSE', 'R2 with CV'])
```

```
metrics = metrics.append(metrics_bp, ignore_index=True)
```

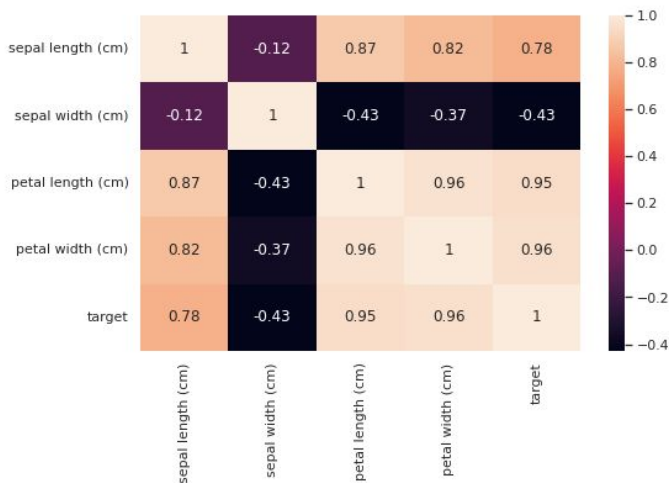
```
metrics.head()
```

	Model	R2	MSE	R2 with CV
0	Linear Regression	0.255382	3761.752362	0.482316
1	Support Vector Machines	0.126417	4413.269770	0.161335
2	SVR best params	0.273890	3668.252470	0.474808

Unsupervised

```
corr = df.corr()  
sns.heatmap(corr, annot=True)
```

<AxesSubplot:>



```
from sklearn import neighbors, datasets, preprocessing  
from sklearn.model_selection import train_test_split  
from sklearn.metrics import accuracy_score
```

```
iris = datasets.load_iris()
```

```
import pandas as pd  
import numpy as np
```

```
df = pd.DataFrame(data= np.c_[iris['data'], iris['target']],  
                  columns= iris['feature_names'] + ['target'])  
df.head()
```


	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	target
0	5.1	3.5	1.4	0.2	0.0
1	4.9	3.0	1.4	0.2	0.0
2	4.7	3.2	1.3	0.2	0.0
3	4.6	3.1	1.5	0.2	0.0
4	5.0	3.6	1.4	0.2	0.0

3 PCA (Principal Component Analysis)

```
from sklearn.decomposition import PCA
```

```
X = df.drop(['target'],axis=1)  
y = df['target']
```

```
pca = PCA().fit(X)  
plt.plot(np.cumsum(pca.explained_variance_ratio_))  
plt.xlabel('number of components')  
plt.ylabel('cumulative explained variance');
```



```
pca = PCA(n_components=1)  
pca.fit(X)  
df_pca = pd.DataFrame(pca.transform(X), columns=['pca'], index=df.index)  
df_pca.head()
```

	pca
0	1.122642
1	-0.000932
2	0.454432
3	0.225439
4	1.349071

4 Models

```

from sklearn.cluster import KMeans
from kneed import KneeLocator


def elbow_plot(df):
    """Create elbow plot from normalized data"""
    sse = {}
    sse_r = []
    for k in range(1, 21):
        kmeans = KMeans(n_clusters=k, random_state=1)
        kmeans.fit(df)
        sse[k] = kmeans.inertia_
        sse_r.append(kmeans.inertia_)

    plt.title('Elbow plot for K selection')
    plt.xlabel('k')
    plt.ylabel('SSE')
    sns.pointplot(x=list(sse.keys()),
                  y=list(sse.values()))

    plt.show()
    return sse_r

```

```
sse = elbow_plot(df_pca)
```



```

kl = KneeLocator(range(1, 21), sse, curve="convex", direction="decreasing")
kl.elbow

```

4

```

k_means = KMeans(n_clusters=4, random_state=42)
k_means.fit(X)

```

...

```

y_pred = k_means.predict(X)
y_pred

```

```

array([[1, 1, 1, 1, 3, 3, 1, 1, 1, 1, 3, 1, 1, 1, 3, 3, 3, 1, 3, 3, 1, 3,
        3, 1, 1, 1, 1, 1, 1, 1, 1, 1, 3, 3, 1, 1, 1, 3, 1, 1, 1, 2, 1, 1,
        3, 1, 3, 1, 3, 1, 0, 0, 0, 2, 0, 0, 0, 2, 0, 2, 2, 0, 2, 0, 0, 0,
        0, 2, 2, 2, 0, 0, 2, 0, 0, 0, 0, 0, 0, 2, 2, 2, 2, 2, 0, 1, 0, 2,
        0, 2, 2, 0, 2, 2, 2, 0, 0, 0, 2, 0, 0, 2, 0, 0, 0, 0, 2, 0, 2, 3,
        0, 2, 0, 2, 0, 0, 0, 3, 2, 2, 0, 0, 0, 2, 0, 0, 0, 0, 0, 0, 0, 3,
        0, 0, 2, 0, 0, 0, 0, 0, 0, 2, 0, 0, 2, 0, 0, 0, 0], dtype=int32)

```





```

df['Cluster'] = y_pred
df.head()

```


	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	target	Cluster
0	0.222222	1.019004	0.067797	0.041667	0	1
1	0.166667	-0.131979	0.067797	0.041667	0	1
2	0.111111	0.328414	0.050847	0.041667	0	1
3	0.083333	0.098217	0.084746	0.041667	0	1
4	0.194444	1.249201	0.067797	0.041667	0	3

Unsupervised




Unsupervised

```
ax = sns.countplot(x="target", hue="Cluster", data=df, palette="Set2")
```



```
ax = sns.countplot(x="Cluster", hue="target", data=df, palette="Set3")
```



Ja hem acabat la part 2!

Gràcies a tots!

